

Ceph SPDK NVMe-oF Gateway Evaluation on openEuler

Xinliang Liu, Senior Engineer, Server, Linaro – sig-SDS
openEuler sig-SDS Meetup 2024 October

Agenda


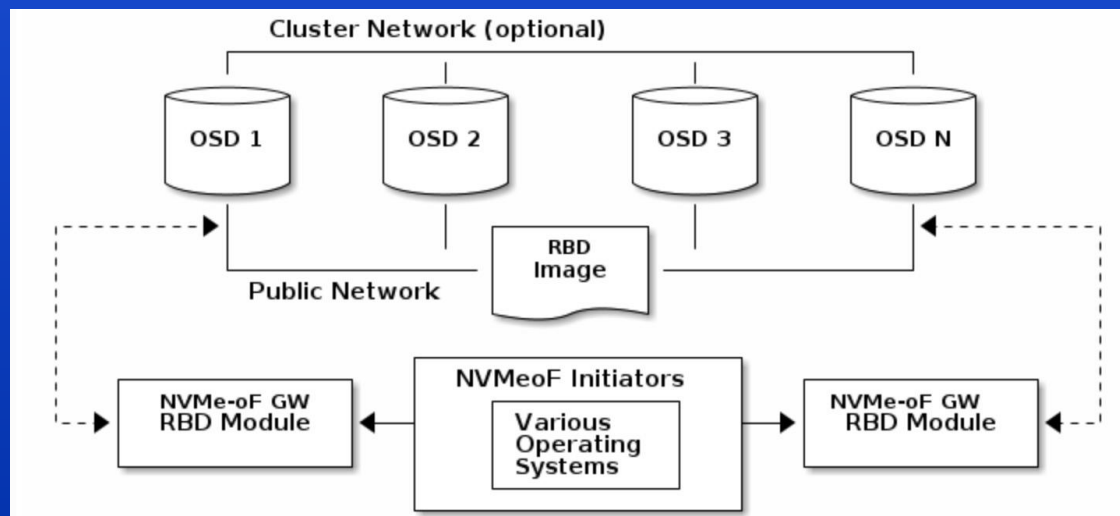
- Ceph SPDK NVMe-oF Gateway介绍
- Ceph SPDK NVMe-oF Gateway Arm上游支持
- Ceph SPDK NVMe-oF Gateway性能测试

Ceph SPDK NVMe-oF Gateway介绍



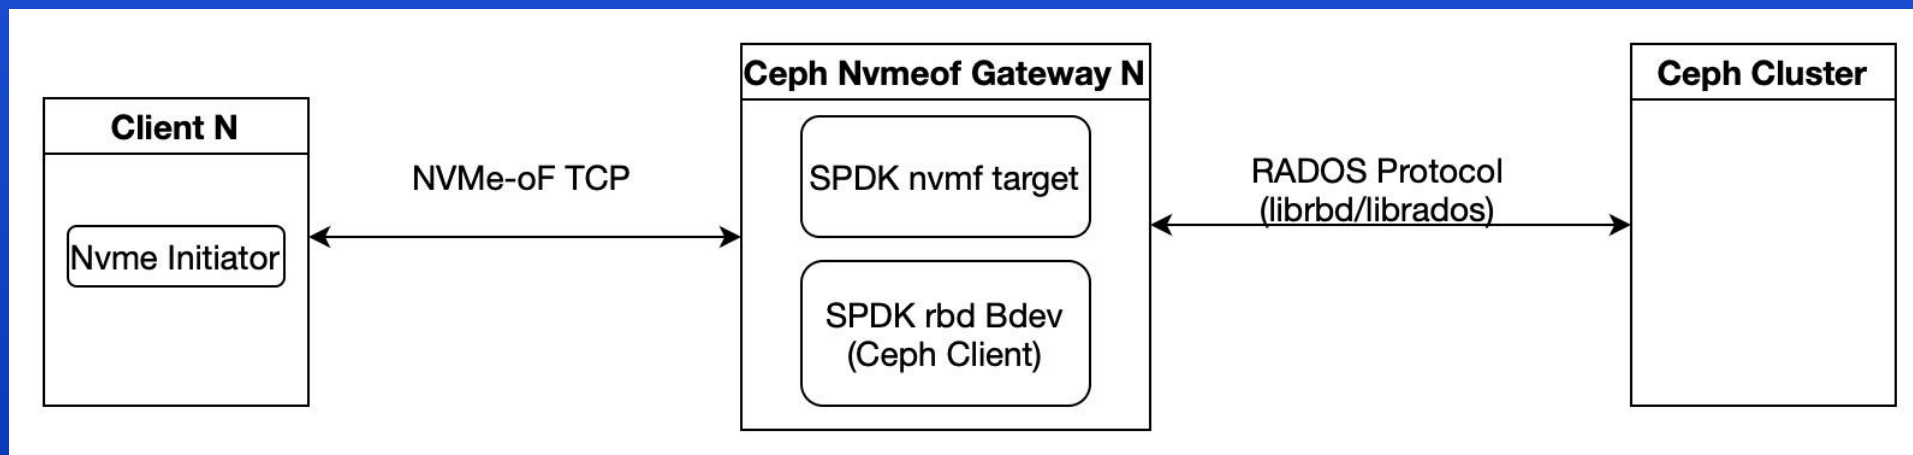
Ceph SPDK NVMe-oF Gateway介绍 - Overview

- What and why is NVMe-oF Gateway?
 - Replace Ceph-iSCSI, convert RBD to NVMe-of Disk gateway
 - NVMe-oF SAN similar network disk solution
 - Very low client %CPU, move ceph client to gw node, more generic network storage protocol.
- Usecase
 - Possibly all the SAN and NVMe-oF usecases, requires client support nvme initiator
 - For no ceph client support environment
 - VMware Virtual Machine, bare metal boot volumes, Virtual Desktop
 - PDUs and hardware accelerators
 - Databases and Analytics, AI etc.
- References
 - <https://www.spiceworks.com/tech/networking/articles/storage-area-network/>
 - <https://www.starwindsoftware.com/blog/what-is-nvme-of-nvme-over-fabrics/>



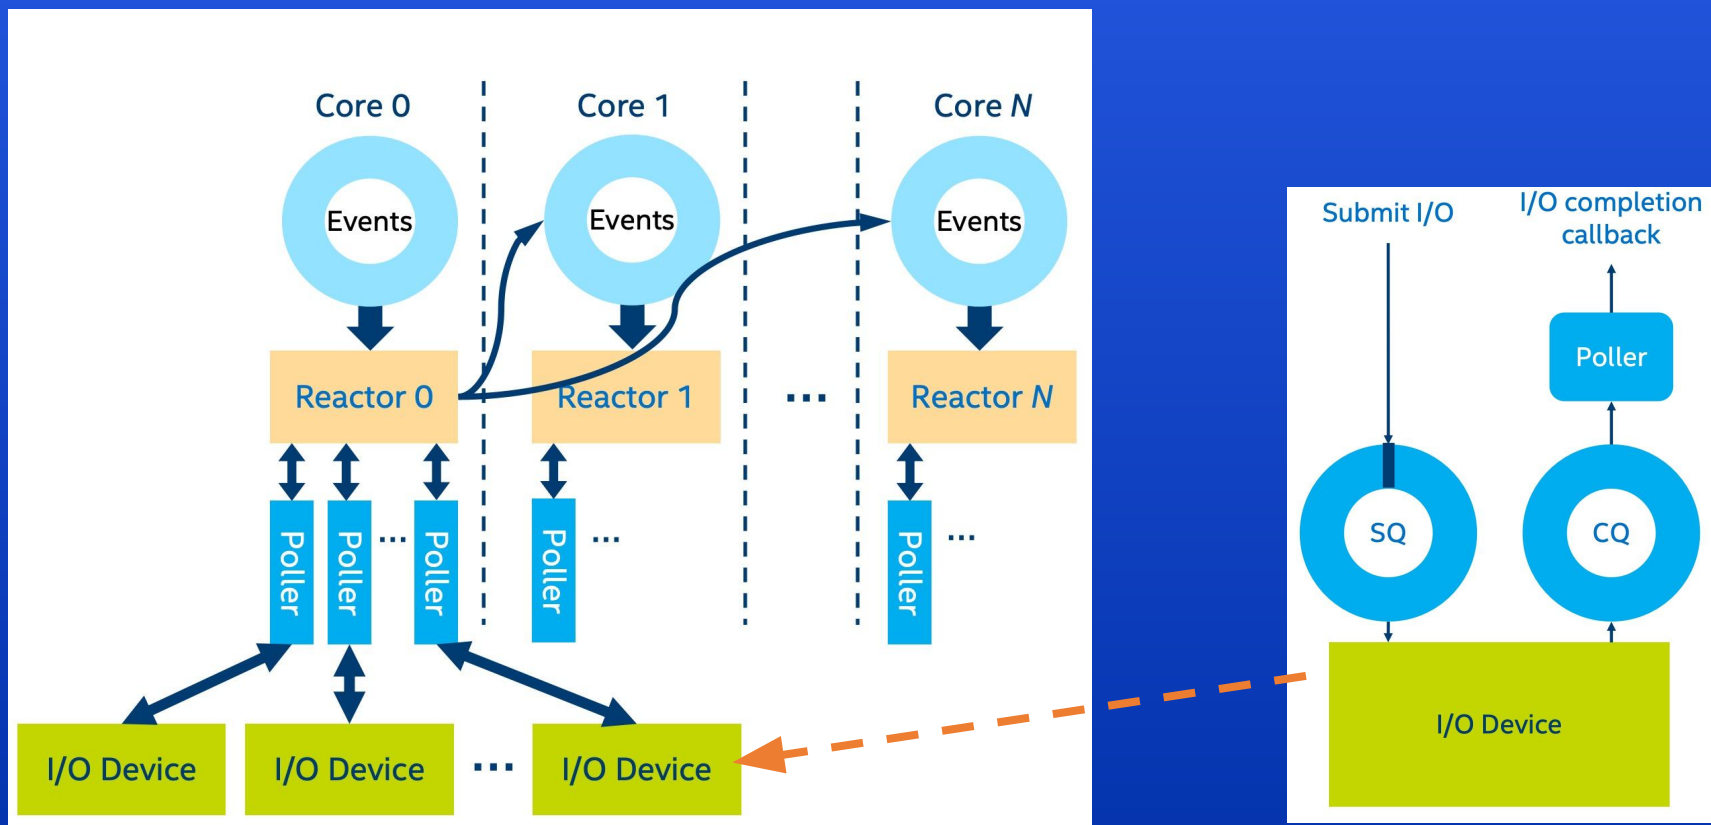
Ceph SPDK NVMe-oF Gateway介绍 - Overview

- The NVMe-oF Gateway presents an NVMe-oF target that exports RADOS Block Device (RBD) images as NVMe namespaces.
- The NVMe-oF protocol allows clients (initiators) to send NVMe commands to storage devices (targets) over a TCP/IP network.
- Enabling clients without native Ceph client support to access Ceph block storage.
- More details see: <https://docs.ceph.com/en/latest/rbd/nvmeof-overview/>



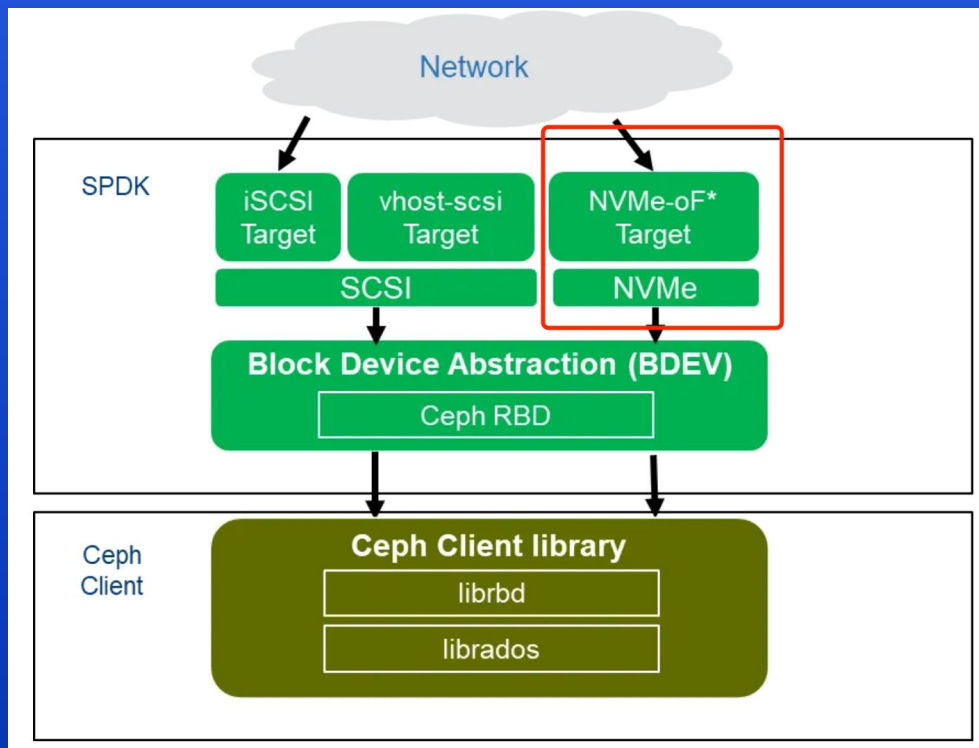
Ceph SPDK NVMe-oF Gateway介绍 - Overview

- SPDK overview
- References:
 - <https://spdk.io/doc/overview.html>
 - <https://www.cnblogs.com/yi-mu-xi/p/12821103.html>



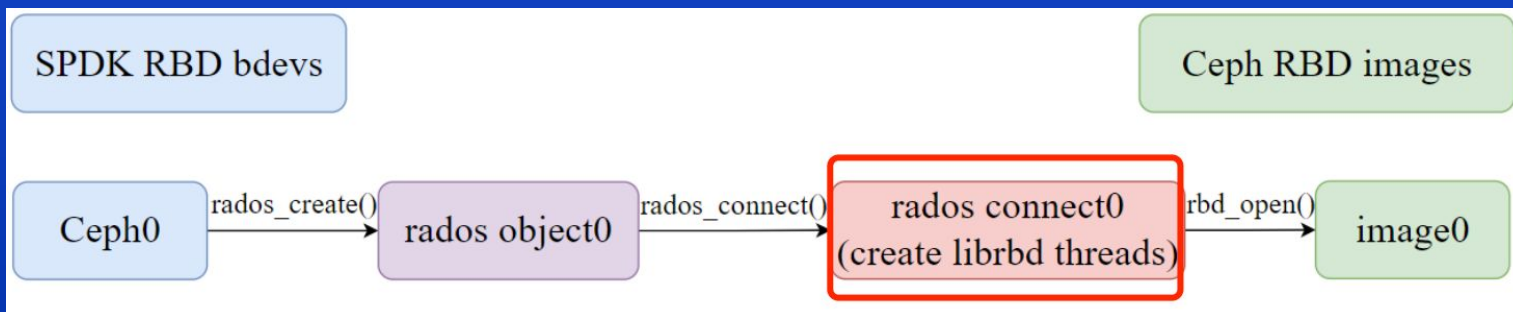
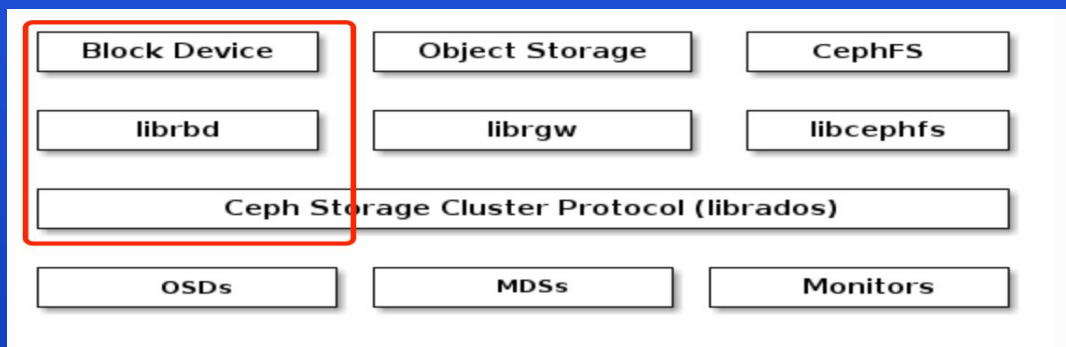
Ceph SPDK NVMe-oF Gateway介绍 - SPDK RBD Bdev

- Convert Ceph RBD to SPDK bdev.
- Allow export to iSCSI/vhost-user/NVMe-oF targets.
- Reference: [管中窥豹SPDK RBD bdev 模块](#)



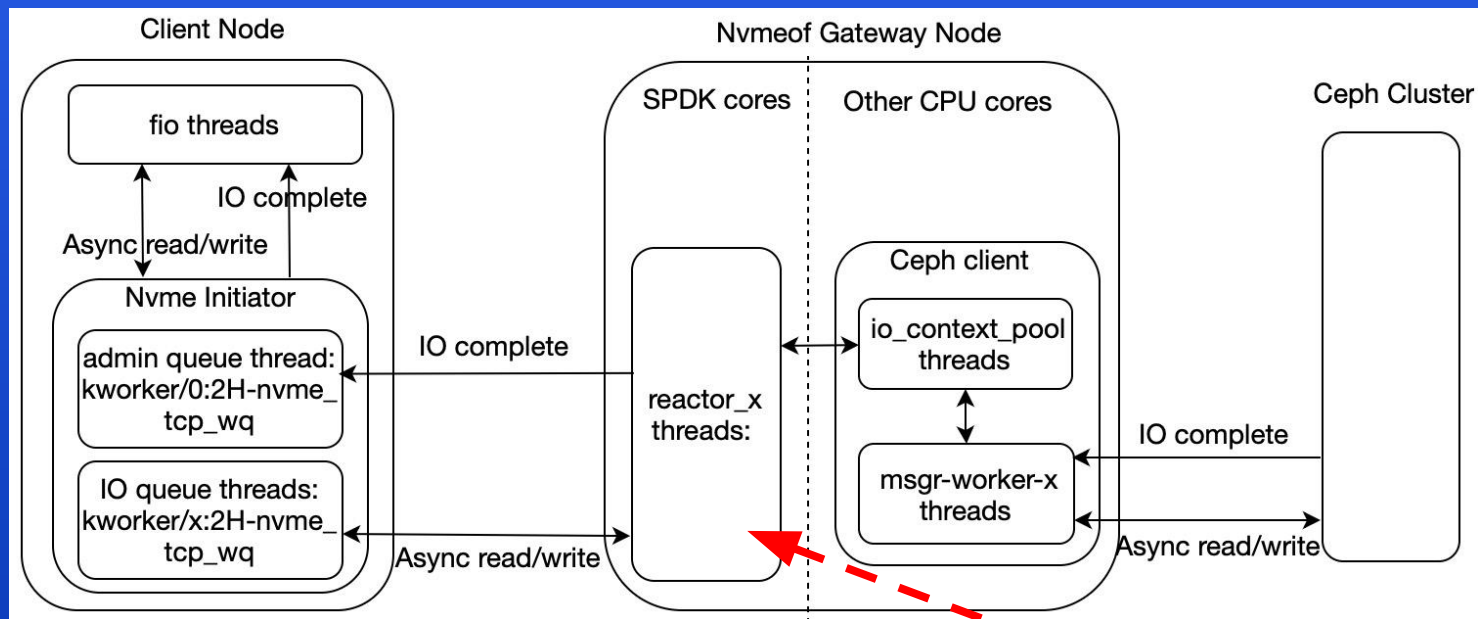
Ceph SPDK NVMe-oF Gateway介绍 - Ceph rbd client

1. Rados cluster object create: `rados_create()`
2. Connect to cluster: `rados_connect()`
 - Create `io_context_pool` threads: io prepare and put to msg queues
 - Create `msggr-worker-x` threads: msg send/recv
3. Rbd io/image ctx open: `rados_ioctx_create()`, `rbd_open()`
4. Rbd image async read/write: `rbd_aio_read/readv/write/writev/flush()`
5. More details see: https://github.com/spdk/spdk/blob/master/module/bdev/rbd/bdev_rbd.c



Ceph SPDK NVMe-oF Gateway介绍 - thread model

Nvmeof Gateway whole thread model



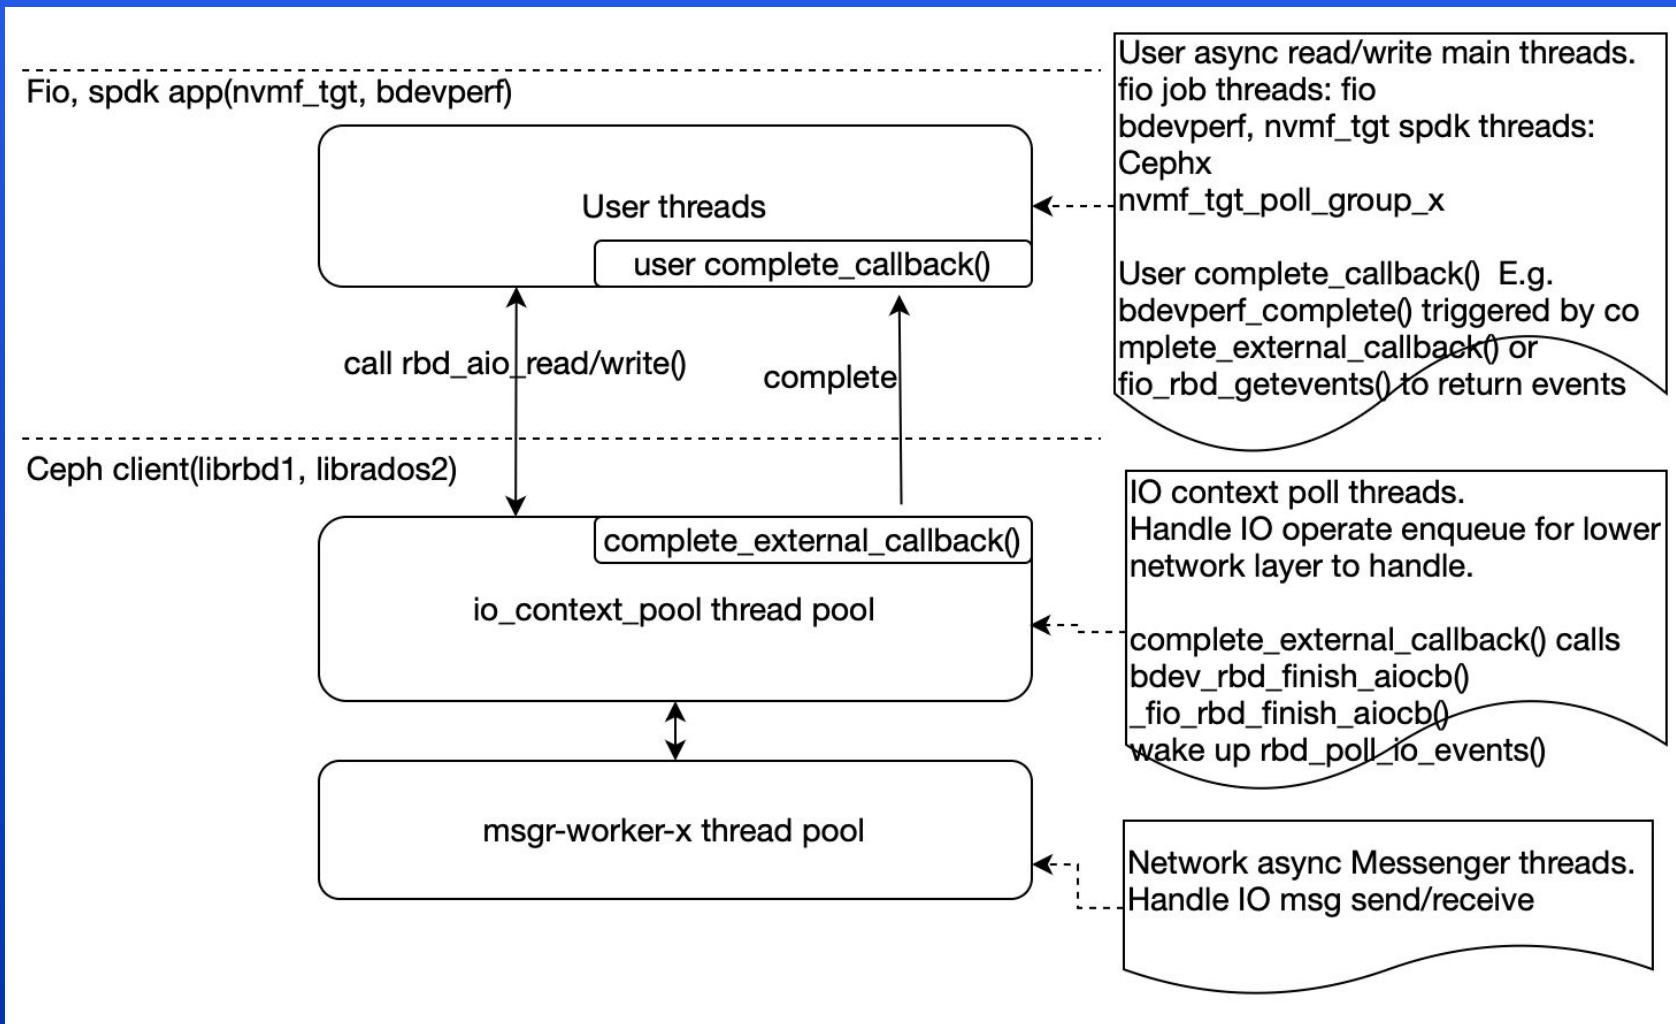
```
reactor_run()
{
    while (1) {
        // run the reactor event ring functoin()
        event_queue_run_batch(reactor);

        // poll thread for each thread nvme_tgt_poll_group_XXX
        spdk_thread_poll() {
            // run the thread msg ring function()
            // e.g. bdevperf_complete()
            msg_queue_run_batch(thread, max_msgs);

            // poll all the pollers of the thread,
            // e.g. nvme_poll_group_poll
            thread_execute_poller()
        }
    }
}
```

Ceph SPDK NVMe-oF Gateway介绍 - thread model

Gateway node thread model



Ceph SPDK NVMe-oF Gateway Arm上游支持



Ceph SPDK NVMe-oF Gateway Arm上游支持 – Enable

Enable building nvmeof container images for Arm64

- Ticket: <https://github.com/ceph/ceph-nvmeof/issues/781>
- PRs
 - Fix x86 hardcode: <https://github.com/ceph/ceph-nvmeof/pull/782>
 - Add Arm build: <https://github.com/ceph/ceph-nvmeof/pull/783>
- Arm Build Note: <https://github.com/ceph/ceph-nvmeof/blob/devel/README.md>

NOTE: For Arm64 build, the default SPDK building SoC is `generic`. To build SPDK for other SoC you need to override the default values of `SPDK_TARGET_ARCH` and `SPDK_MAKEFLAGS`. To know which values to set for all the supported Arm64 SoCs see [the socs and implementer_xxx parts](#). E.g. for kunpeng920 SoC:


```
make build SPDK_TARGET_ARCH="armv8.2-a+crypto" \  
SPDK_MAKEFLAGS="DPDKBUILD_FLAGS=-Dplatform=kunpeng920"
```



Ceph SPDK NVMe-oF Gateway Arm上游支持 – CI

Add build-arm64 job, multi-platform build

- Ticket: <https://github.com/ceph/ceph-nvmeof/issues/805>
- PR: <https://github.com/ceph/ceph-nvmeof/pull/806>
 - Migrate docker-compose from v1 to v2.
 - Support qemu multi-platform container image build.
 - Add build-arm64 job run per PR and daily.
 - Build and publish nvmeof, nvmeof-cli images to quay.io/ceph/.



← ceph/nvmeof 40264b7dfa24

MANIFEST	OPERATING SYSTEM	ARCHITECTURE
SHA256 0a6ef3fe1151	linux	arm64
SHA256 8eaabefa8ee6	linux	amd64

Ceph SPDK NVMe-oF Gateway性能测试



Ceph SPDK NVMe-oF Gateway性能测试 - Testbed

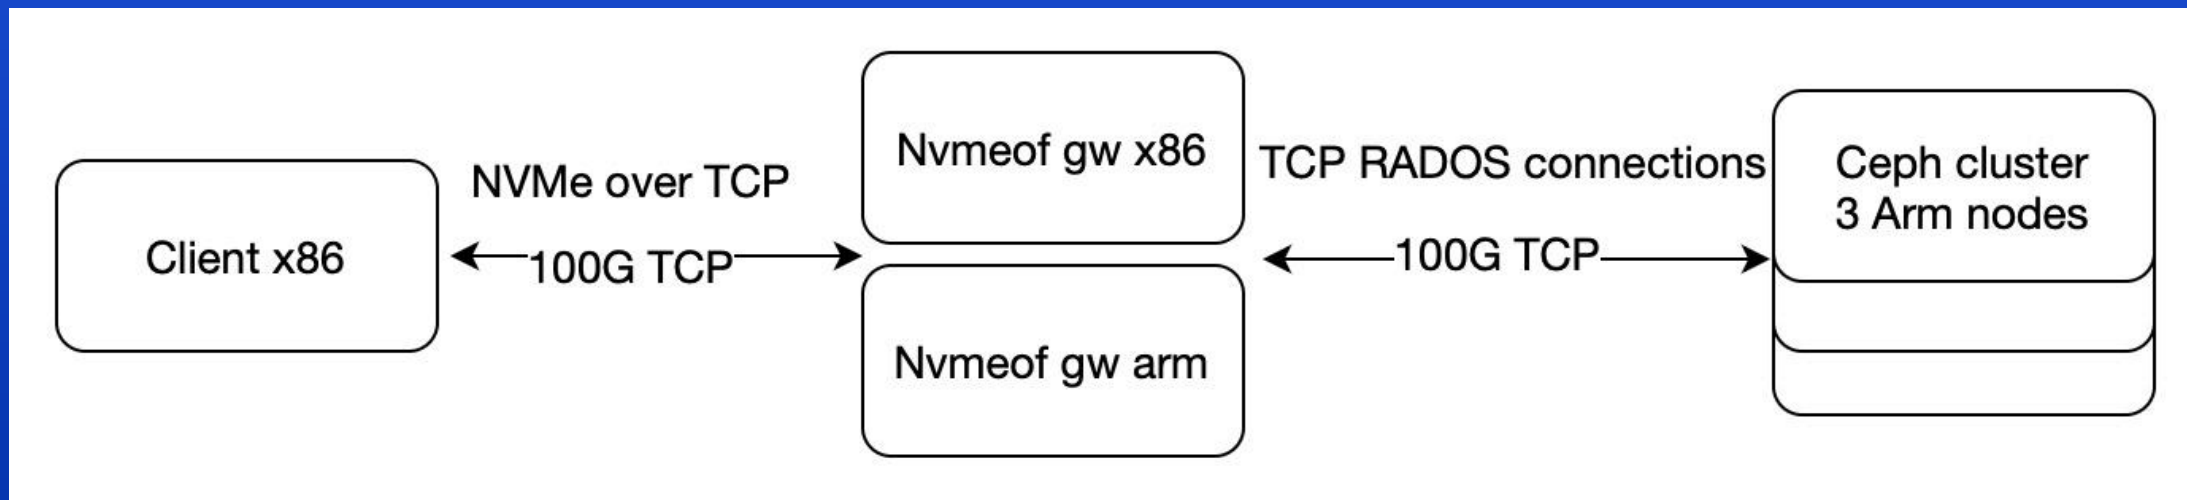
Test Environment

Hardware

- Arm CPU: Kunpeng 920, 2.6GHz, 4 numa nodes
- X86 (Gateway) CPU: Intel(R) Xeon(R) Platinum 8180 CPU @ 2.50GHz, 112 cpus, 2 numa nodes
- x86-2 (Client) CPU: Intel(R) Xeon(R) Gold 6248R CPU @ 3.00GHz, 96 cpus, 2 numa nodes
- Disk: 3 x ES3000 V6 NVMe SSD 3.2T per Arm server
- Network: 1x MLNX ConnectX-5 100Gb IB, 1x1G tcp

Software

- OS: openEuler 22.03 LTS SP3, kernel 5.10.0-192.0.0.105.oe2203sp3
- Ceph: [main-nvmeof](#) require revert commit "nvmeof gw monitor: disable by default" ←
- SPDK: 24.05
- nvmeof: 1.3.2
- fio: fio-3.29



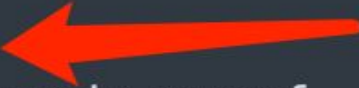
Ceph SPDK NVMe-oF Gateway性能测试 - Testbed

Arm and x86 hybrid deployment

Reference:

- <https://docs.ceph.com/en/latest/cephadm/install/#bootstrap-a-new-cluster>
- <https://docs.ceph.com/en/latest/cephadm/operations/#purging-a-cluster>

```
# Deploy Ceph Cluster, mix arches x86 and arm64
wget -c https://download.ceph.com/rpm-reef/el9/noarch/cephadm
chmod +x cephadm
cephadm --image quay.io/xin3liang0/ceph:main-nvmeof bootstrap --mon-ip 192.168.0.86
cephadm shell
ceph config set mgr mgr/cephadm/use_repo_digest false
ceph config set global container_image quay.io/xin3liang0/ceph:main-nvmeof
cephadm shell
ceph orch host add server5
ceph orch host add server1
# osds_per_device: 8 # Run more osd daemons to utilize a nvme disk
ceph orch apply -i osd_spec.yml
```



Ceph SPDK NVMe-oF Gateway性能测试 - Testbed

Arm and x86 hybrid deployment

Reference

- <https://docs.ceph.com/en/latest/rbd/nvmeof-target-configure/>
- https://ci.spdk.io/download/2022-virtual-forum-prc/D2_4_Yue_A_Performance_Study_for_Ceph_NVMeoF_Gateway.pdf
- https://ci.spdk.io/download/2022-virtual-forum-prc/D2_3_Yifan_The_usage_of_SPDK_RBD_bdev_and_performance_tuning.pdf

```
# Deploy NVMe-oF Gateway service
ceph config set mgr mgr/cephadm/container_image_nvmeof quay.io/xin3liang0/nvmeof:latest
nvmeof_pool=nvmeof
ceph osd pool create $nvmeof_pool
rbd pool init $nvmeof_pool
ceph orch apply nvmeof $nvmeof_pool default --placement="server3,client10"
vi /var/lib/ceph/23190a0a-4fba-11ef-9b42-60d755fdcf8c/nvmeof.nvmeof.default.server3.hcobql/ceph-nvmeof.conf
# Bind SPDK and ceph client threads in the same numa
[spdk]
tgt_cmd_extra_args = "-m 0xF0000000"
librbd_core_mask = 0xFFFFF000000000
vi /var/lib/ceph/23190a0a-4fba-11ef-9b42-60d755fdcf8c/nvmeof.nvmeof.default.server3.hcobql/config # ceph.conf
# Increase io_context_pool and msgr-worker-x thread number
[global]
ms_async_op_threads = 9
librados_thread_count = 6
ceph orch daemon restart nvmeof.nvmeof.default.client10.nvycr
ceph orch daemon restart nvmeof.nvmeof.default.server3.hcobql
```

Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

FIO config

- FIO rbd config

```
(.venv) [root@client1 spdktest]# cat fio_test-rbd.conf
[global]
#stonewall
description="Run ${RW} ${BS} rbd test"
bs=${BS}
ioengine=rbd
clientname=admin
pool=nvmeof
thread=1
group_reporting=1
direct=1
verify=0
norandommap=1
time_based=1
ramp_time=10s
runtime=1m
iodepth=${IODEPTH}
rw=${RW}
numa_cpu_nodes=0

[test-job1]
rbdname=fio_test_image1
```

Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

FIO config

- FIO nvmeof config

```
(.venv) [root@client1 spdktest]# cat fio_test-nvmeof.conf
[global]
#stonewall
description="Run ${RW} ${BS} NVMe ssd test"
bs=${BS}
#ioengine=libaio
ioengine=io_uring
thread=1
group_reporting=1
direct=1
verify=0
norandommap=1
time_based=1
ramp_time=10s
runtime=1m
iodepth=${IODEPTH}
rw=${RW}
numa_cpu_nodes=0

[test-job1]
#filename=/dev/nvme2n1
filename=/dev/nvme2n2
#test-job2
```

Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

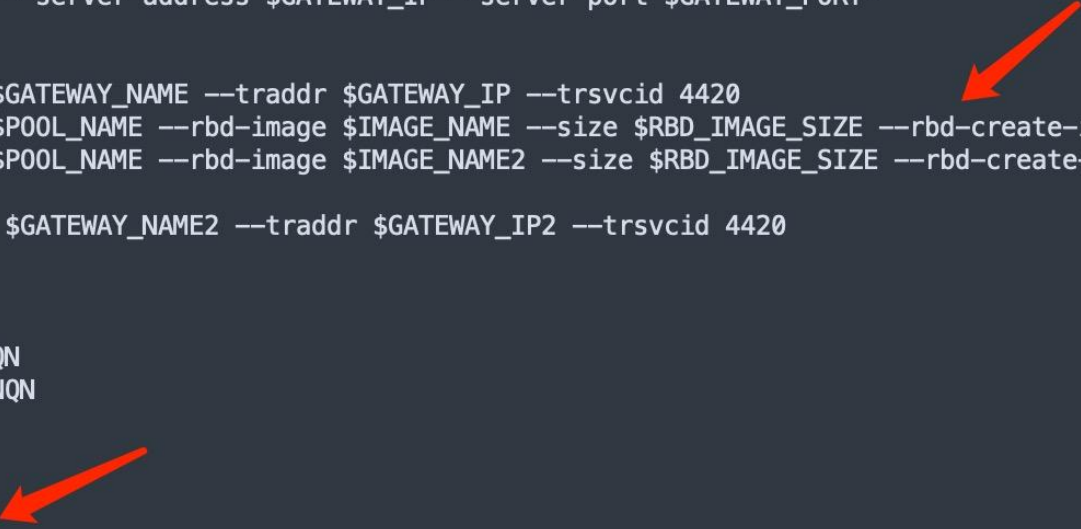
Create test rbd

```
GATEWAY_NAME2=client10
GATEWAY_IP2=192.168.0.100
GATEWAY_NAME=server3
GATEWAY_IP=192.168.0.83
GATEWAY_PORT=5500
SUBSYSTEM_NQN=nqn.2016-06.io.spdk:cnode1
POOL_NAME=nvmeof
IMAGE_NAME=rbd1
IMAGE_NAME2=rbd2
RBD_IMAGE_SIZE=20G
NVMEOF_CLI="podman run --rm quay.io/xin3liang0/nvmeof-cli:latest --server-address $GATEWAY_IP --server-port $GATEWAY_PORT"

$NVMEOF_CLI subsystem add --subsystem $SUBSYSTEM_NQN
$NVMEOF_CLI listener add --subsystem $SUBSYSTEM_NQN --host-name $GATEWAY_NAME --traddr $GATEWAY_IP --trsvcid 4420
$NVMEOF_CLI namespace add --subsystem $SUBSYSTEM_NQN --rbd-pool $POOL_NAME --rbd-image $IMAGE_NAME --size $RBD_IMAGE_SIZE --rbd-create-image
$NVMEOF_CLI namespace add --subsystem $SUBSYSTEM_NQN --rbd-pool $POOL_NAME --rbd-image $IMAGE_NAME2 --size $RBD_IMAGE_SIZE --rbd-create-image
$NVMEOF_CLI host add --subsystem $SUBSYSTEM_NQN --host "*"
$NVMEOF_CLI2 listener add --subsystem $SUBSYSTEM_NQN --host-name $GATEWAY_NAME2 --traddr $GATEWAY_IP2 --trsvcid 4420

dnf install nvme-cli
modprobe nvme-fabrics
nvme connect -t tcp --traddr $GATEWAY_IP -s 4420 -n $SUBSYSTEM_NQN
nvme connect -t tcp --traddr $GATEWAY_IP2 -s 4420 -n $SUBSYSTEM_NQN
nvme list

rbd create --size $RBD_IMAGE_SIZE $nvmeof_pool/fio_test_image1
rbd create --size $RBD_IMAGE_SIZE $nvmeof_pool/fio_test_image2
```




Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

SSD rbd Preconditioning(预处理)

- Sequential write 2X+ size of rbd to get SSD steady state.
 - RW=write BS=128k IODEPTH=128 fio ./[fio_test-rbd.conf|fio_test-nvmeof.conf] --numjobs=1 # run 1 min
- Reference
 - SSD Preconditioning: https://ci.spdk.io/download/performance-reports/SPDK_nvme_bdev_perf_report_2405.pdf

```
[ceph: root@server6 ~]# rbd -p nvmeof disk-usage
NAME                PROVISIONED  USED
fio_test_image1     20 GiB      20 GiB
fio_test_image2     20 GiB       0 B
rbd1                 20 GiB      20 GiB
rbd2                 20 GiB      20 GiB
<TOTAL>             80 GiB      60 GiB
[ceph: root@server6 ~]#
```

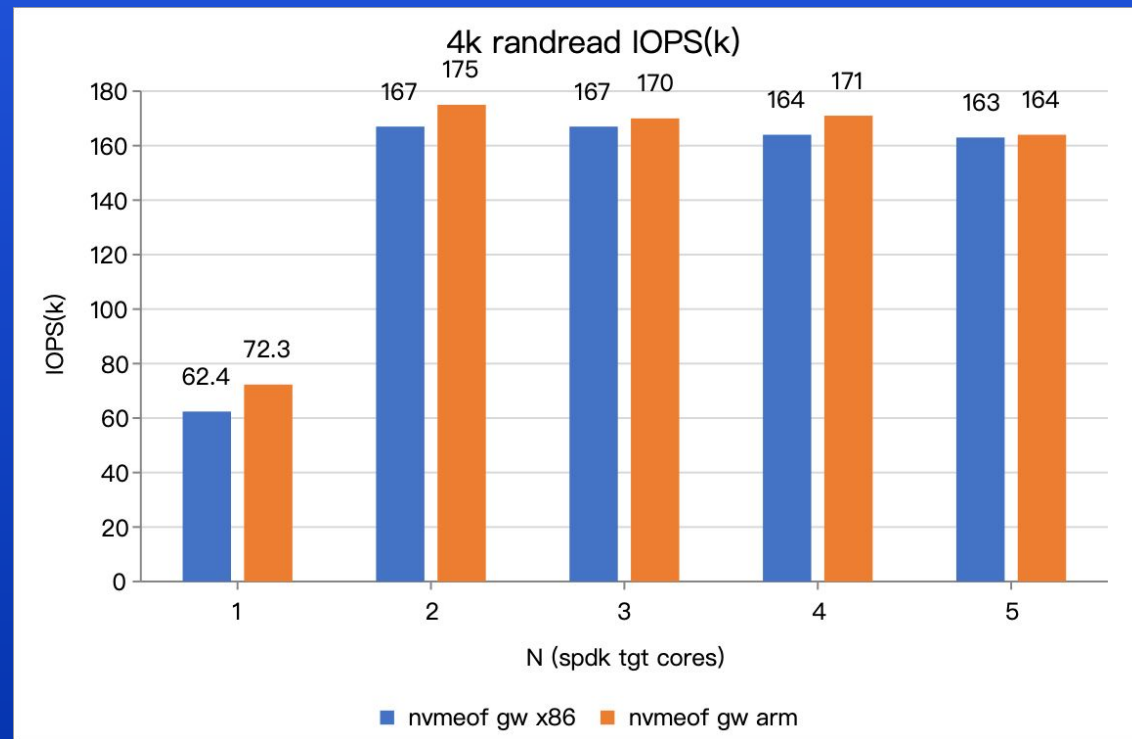
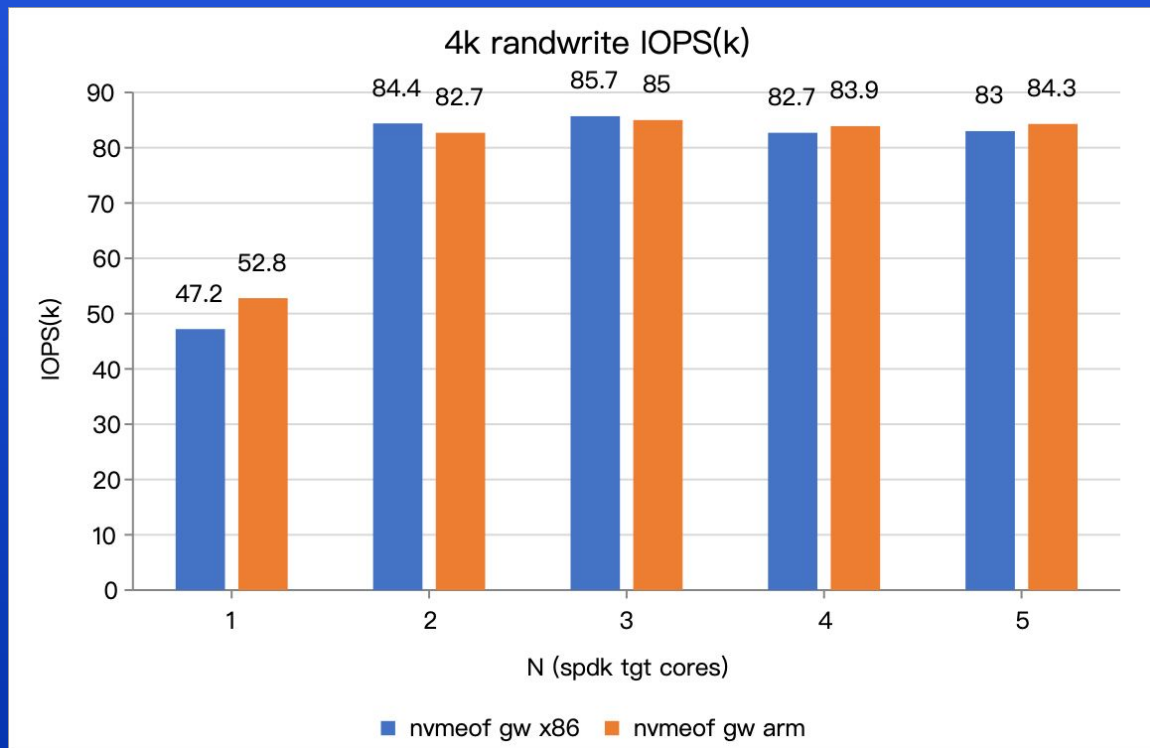


Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

Case 1: scale spdk cores (triple ceph client msg and io enqueue threads, bind ceph client, spdk tgt threads in the same numa as NIC)

4k RW=randwrite|randread IODEPTH=128 fio ~/spdktest/fio_test-nvmeof.conf --numjobs=8

- ~50k IOPS per core
- Find out cpus via spdk_top cpu not 100%


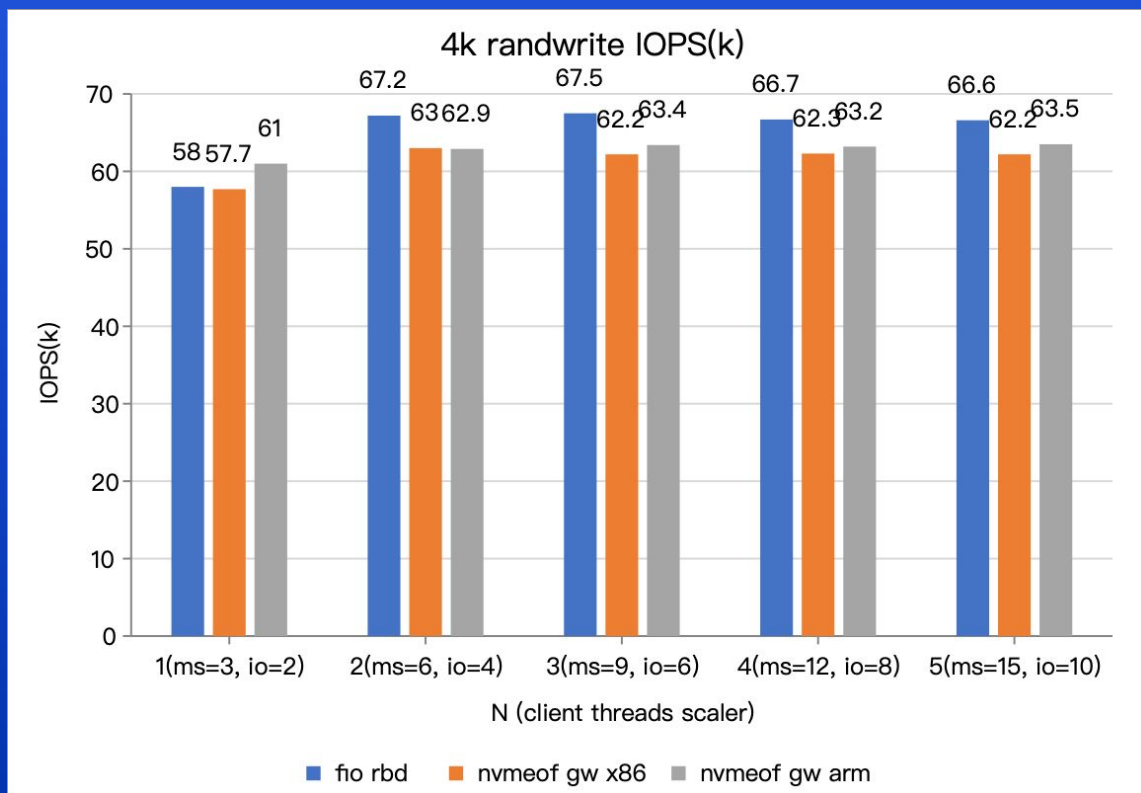


Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

Case 2: scale ceph client msg and io enqueue threads (4 spdk cores , bind ceph client, spdk tgt threads in the same numa as NIC)

4k RW=randwrite IODEPTH=128 fio ~/spdktest/[fio_test-rbd.conf|fio_test-nvmeof.conf] --numjobs=1

- ms_async_op_threads = 3 * N , librados_thread_count = 2 * N
- Find out thread's cpu not 100% via top, N = 2 fio (N = 3 nvmeof) get best performance number



Ceph SPDK NVMe-oF Gateway性能测试 - 1 rbd test

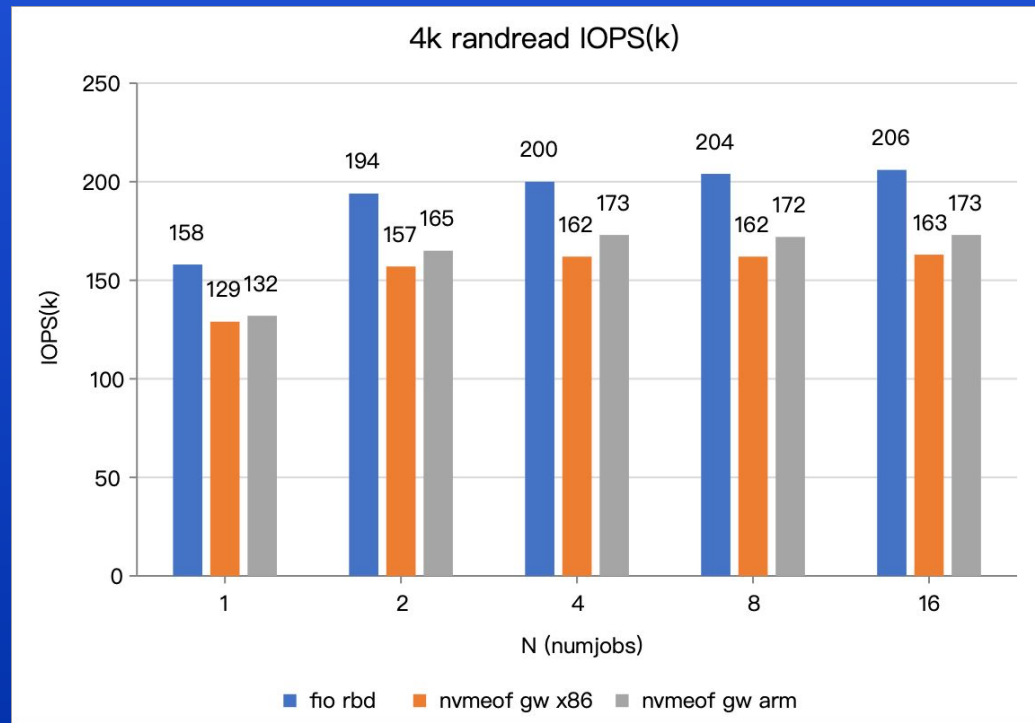
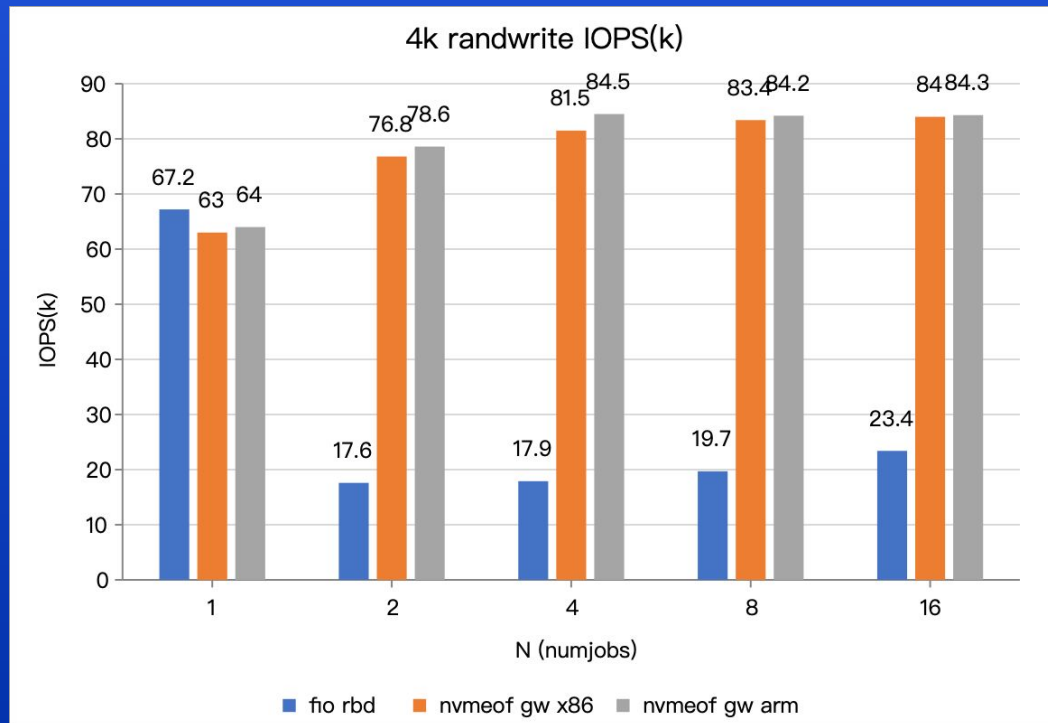
Case 3: double/triple client msg and io enqueue threads (4 spdk cores, bind ceph client, spdk tgt threads in the same numa as NIC)

4k RW=randwrite IODEPTH=128 fio ~/spdktest/[fio_test-rbd.conf|fio_test-nvmeof.conf] --numjobs=N

- FIO rbd: ms_async_op_threads = 6 # 3->6, librados_thread_count = 4 # 2->4
- FIO rbd: Should numjobs=1, because 1 rados connection, per job multi rados connect performance drop

Due to the rbd exclusive write lock


- Nvmeof : ms_async_op_threads = 9 # 3->9 , librados_thread_count = 6 # 2->6 , X86 gw ~20% drop, Arm gw ~15% drop



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test

Case 1: fio rbd multi-client test for ceph cluster capability (numjob=1, client_num=6)

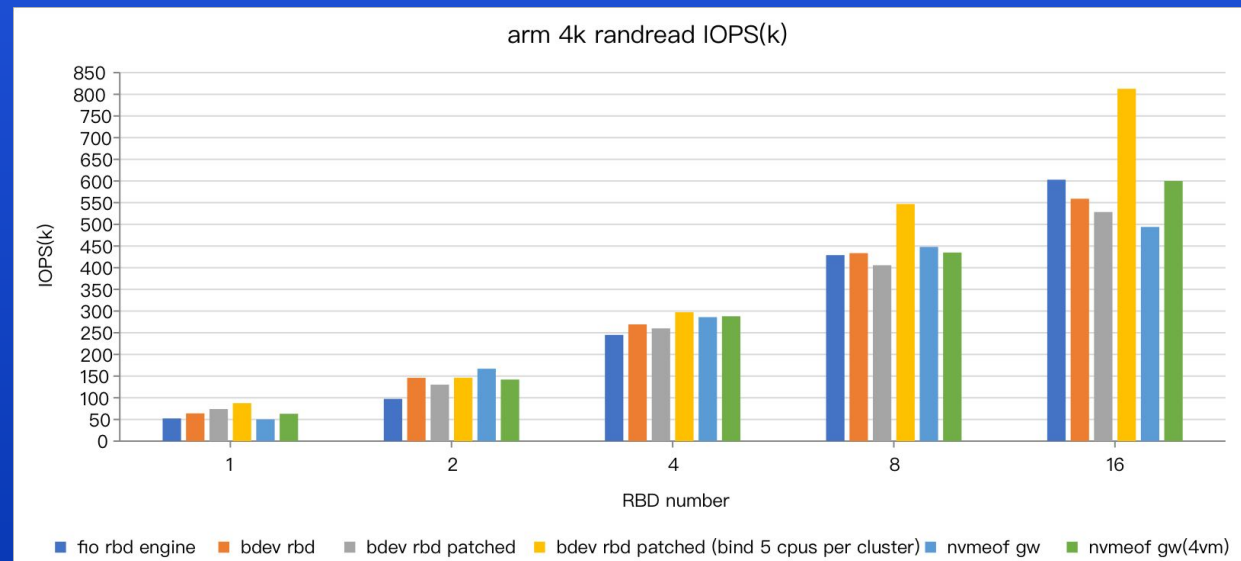
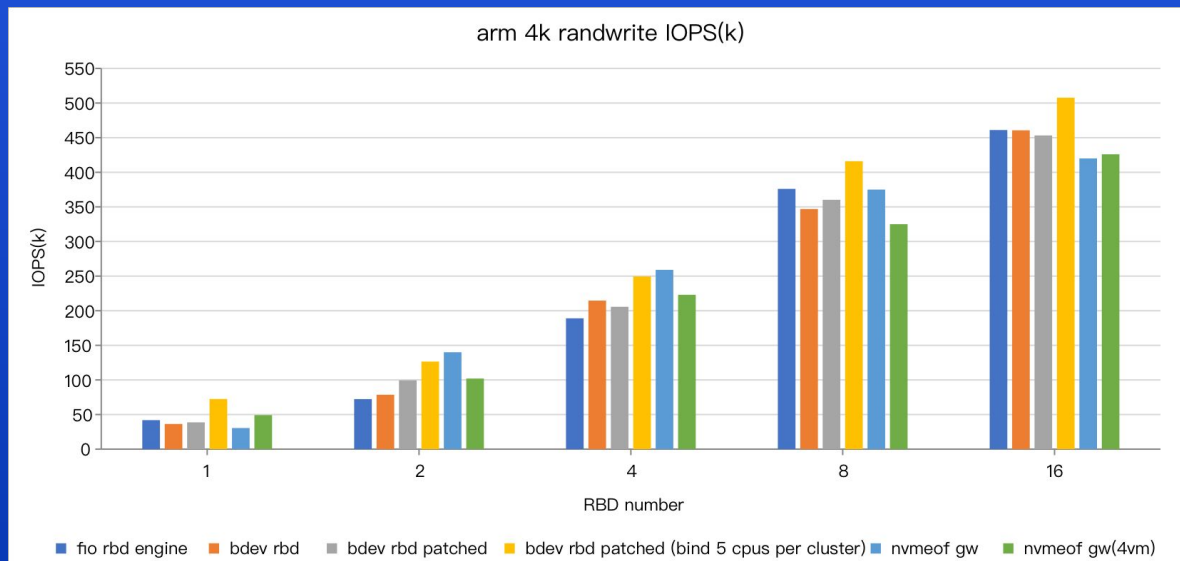
- `./fio -server& #` run in the clients nodes
- `./run-n-fio-rbd-test_multi-client.sh 48 #` run in one client node



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test


Case 2: fio, bdev, gw rbd test (numjob=1, 16 spdk cores, bdevs_per_cluster = 1, iommu.passthrough=1)

- bdevperf -q 128 -o 4096 -w randwrite -t 60 -m [0-15] -z # run in one terminal
- [./run-n-bdev-rbd-test1.sh](#) N, [./run-n-fio-nvmeof-test1.sh](#) N, [run-n-fio-rbd-test1.sh](#) N,
 - [./run-bdev-rbd-test2.sh](#) N 16-95 # each ceph client threads binding to 5 cpu cores
- Test scripts: <https://github.com/xin3liang/home-bin/tree/master/spdk-fio>



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test tips

- Tip1: Ceph pool pg_num impacts performance much
 - pg_autoscale_mode scales pg_num according pool space utilization
 - Setting pg_num can get the maximum performance at the pool creation beginning
 - `pool=nvmeof; ceph osd pool create $pool --size 1 --pg_num 16384 --pgp_num 16384 --bulk`
 - `#ceph osd pool set $pool target_size_ratio 0.9`
 - `ceph osd pool set $pool pg_autoscale_mode off`
 - `rbd pool init $pool`
 - References
 - [Ceph: A Journey to 1 TiB/s](#)
 - [PG Calc — Ceph Documentation](#)
 - [Placement Groups — Ceph Documentation](#)



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test tips

- Tip2: Nvmeof-gateway should build release type
 - Upstream tacking card: <https://github.com/ceph/ceph-nvmeof/issues/938>
 - SPDK_CONFIGURE_ARGS="--with-rbd --disable-tests --disable-unit-tests --disable-examples --enable-debug"
 - Debug build type performance is pool.
- Tip3: Build LSE atomic instructions for Arm64
 - Check LSE: perf top see __aarch64_xxx_xxx_rel, e.g. __aarch64_ldadd8_acq_rel
 - Check LSE: objdump and egrep -i
'cas|caspl|swp|ldadd|stadd|ldclr|stclr|ldeor|steor|ldset|stset|ldsmx|stsmx|ldsmn|stsmn|ldumx|stumx|ldumn|stumn'
 - Check load exclusives and store exclusives: objdump and egrep -i 'ldxr|ldaxr|stxr|stlaxr'
 - Refer to: <https://learn.arm.com/learning-paths/servers-and-cloud-computing/lse/example/>
- Tip4: Passthrough iommu
 - iommu overhead is nonnegligible
 - Disable or Passthrough it
 - kernel cmdline: iommu.passthrough=1 (arm64), iommu=pt(Intel)

Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test tips

- Tip5: Nvmeof controller IO queue number \geq SPDK CPU cores number (but not too larger than)
 - ceph-nvmeof.conf: "max_io_qpairs_per_ctrlr": 7
 - To utilize all the spdk cores and get best performance.
- Tip6: Scale spdk cores
 - Add spdk cores when spdk_top cpu% is full
 - ceph-nvmeof.conf: tgt_cmd_extra_args = "-m 0xFF"
 - Might need to increase iobuf pool
 - `RPC=/usr/libexec/spdk/scripts/rpc.py`
 - `bdevperf -q 128 -o 4096 -w randwrite -t 60 -m [0-15] -z --wait-for-rpc`
 - `$RPC iobuf_set_options --small-pool-count 32767 --large-pool-count 8191 && $RPC framework_start_init`

[1] THREADS

[2] POLLERS

[3] CORES

Core	Threads	Pollers	Idle [us]	Busy [us]	Busy %	Status	Intr	Sys %	Irq %	CPU %	Freq [MHz]
0	4	5	2110	1172571	99.82	Busy	No	32.95	17.76	99.82	N/A
1	4	4	0	1174489	100.00	Busy	No	31.11	15.09	100.00	N/A
2	3	3	45	1174293	100.00	Busy	No	32.97	17.27	100.00	N/A
3	3	3	3	1174163	100.00	Busy	No	35.56	15.89	100.00	N/A
4	3	3	15	1173860	100.00	Busy	No	34.48	18.69	100.00	N/A

Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test tips

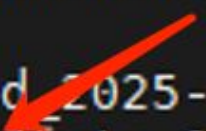
- Tip7: Tune Ceph client thread number
 - ceph-nvmeof.conf: : bdevs_per_cluster = 1 as fio rbd engine
 - <https://github.com/ceph/ceph-nvmeof/issues/939>
 - ceph.conf
 - ms_async_op_threads 3 is default, msgr-worker-x thread number
 - librados_thread_count 2 is default, io_context_pool thread number
 - Tune if top cpu% is near 100%
- Tip8: Bind Ceph client threads
 - Binding cpu could get better performance
 - ceph-nvmeof.conf: librbd_core_mask = 0xFFFF0000
 - bdevperf test script [run-n-bdev-rbd-test2.sh](#)
 - \$RPC bdev_rbd_register_cluster \$cluster --core-mask "\$cpu_list"
 - cpus_per_cluster=5

Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test issue

Issue1: Fio rbd engine can only run small number of jobs

- Upstream tracking issue: <https://github.com/axboe/fio/issues/1778>
 - When running more jobs it will get stuck and no any fio test is running
 - Workaround: running in a container environment without this issue
 - `podman run -it --pids-limit 4096 --name bdevtest centos-stream9`

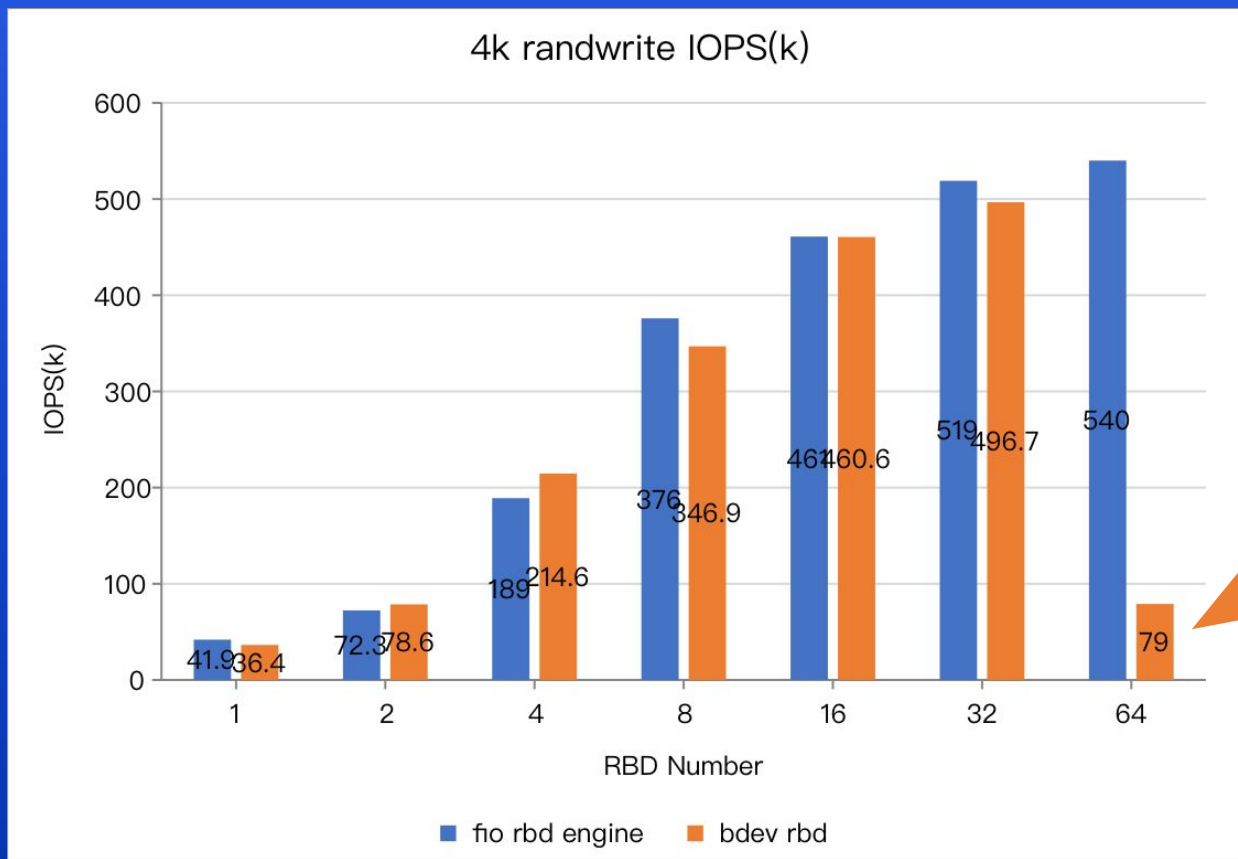
```
[root@server2 spdktest]# ./run-n-fio-rbd-test1.sh 16 randread
fio-3.38-13-gf241
Begin to Run test tmp/fio-rbd_2025-01-21_09-52.log
Jobs: 16 (f=0): [/(16)][-.-%][eta 00m:41s]
```



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test issue

Issue2: 4k randwrite/randread IOPS drops rapidly as the RBD number grows large

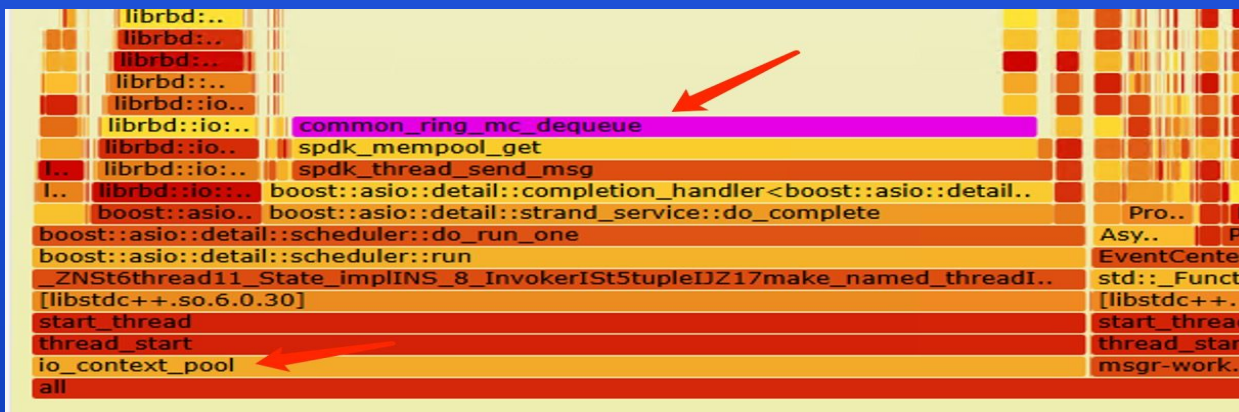
- Upstream tracking issue: <https://github.com/spdk/spdk/issues/3547>



Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test issue

Issue2: 4k randwrite/randread IOPS drops rapidly as the RBD number grows large

- Bottleneck analysis
 - The perf top and flame graph show that common_ring_mc_dequeue() calls spend too much time.
 - io_context_pool threads are non-EAL spdk threads.
 - Due to the lack of per-core caching, rte_mempool_get() performance will suffer when called by unregistered non-EAL threads.



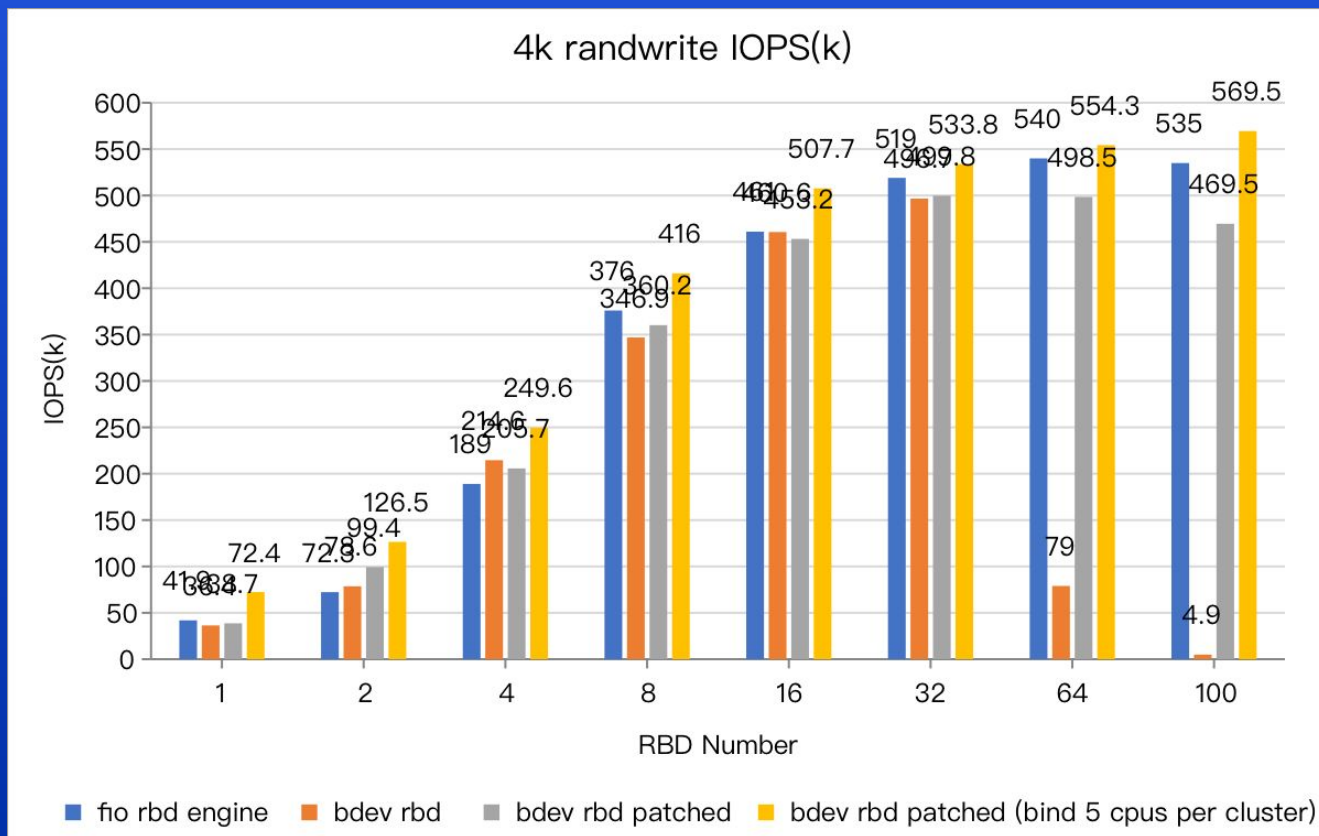
Samples: 1M of event 'cycles:P', 4000 Hz, Event count (approx.): 827427613819 lost: 0/0 drop: 0/0			
Overhead	Shared Object	Symbol	
68.16%	bdevperf	[.] common_ring_mc_dequeue	
5.14%	bdevperf	[.] spdk_ring_dequeue	
2.17%	bdevperf	[.] spdk_thread_poll	
1.91%	[kernel]	[k] finish_task_switch.isra.0	
1.86%	bdevperf	[.] reactor_run	
1.70%	bdevperf	[.] thread_poll	
1.25%	bdevperf	[.] msg_queue_run_batch	
0.88%	libceph-common.so.2	[.] __aarch64_ldadd8_acq_rel	
0.66%	libc.so.6	[.] __int_free	
0.63%	[kernel]	[k] handle_softirqs	

Ceph SPDK NVMe-oF Gateway性能测试 - multi-rbd test issue

Issue2: 4k randwrite/randread IOPS drops rapidly as the RBD number grows large

Possible fix ways:

1. Revert commit "[62210eff55bd](#) bdev/rbd: Remove epoll based group polling mechanism."
2. Or, don't use msg mempool here, just use calloc() to allocate msg
 - [rbd-bdev-drop-rapidly-fix.patch](#) seems work, still under discussing
 - Need to handle msg release properly in msg_queue_run_batch()



Ceph SPDK NVMe-oF Gateway性能测试 - Summary

Summary

- Ceph main requires revert commit "nvmeof gw monitor: disable by default".
- Ceph nvmeof gateway requires rebuilding to non-debug release.
- Build LSE atomic instructions for Arm64.
- RBD Preconditioning: sequential write 2X+ size of rbd to get SSD steady state.
- Add spdk cores when spdk_top cpu% is full
- Bind and tune Ceph client, spdk tgt threads, bdevs_per_cluster=1.
- DPDK mempool is not working well for many non-EAL threads.
- Useful Links
 - <https://spdk.io/cn/articles/>
 - <https://spdk.io/doc/>
 - <https://www.youtube.com/@storageperformancedevelopm301>

THANKS